

Introduction to Reinforcement Learning

Machine Learning

András Attila Sulyok

Pázmány Péter Catholic University
Faculty of Information Technology and Bionics

21/05/2024



What is the goal of this lecture?

- Understand the formulation in RL
 - When can it be used?
 - Why is it difficult?
- Understand the idea behind REINFORCE and Q-learning
 - Two different approaches

Fun Time

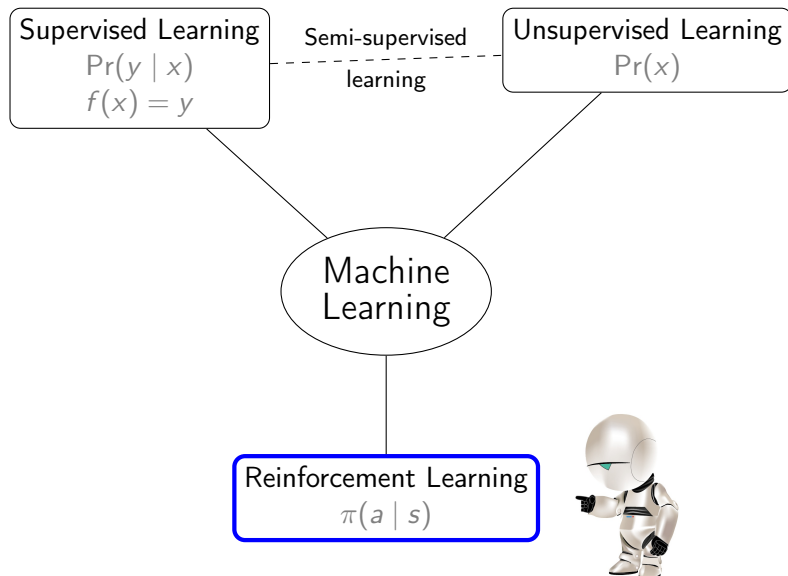
How familiar are you with RL?

- ① With the what?
- ② I know about the theory.
- ③ I have implemented some RL algorithms.
- ④ I contribute new theory to the community.
- ⑤ I have my own army of cyber-lizardpeople AGIs ready to take on the world.


Reference answer: None this time.

- 1 Motivation
- 2 Multi-armed Bandit
- 3 Markov Decision Process
- 4 Policy Gradients
- 5 Temporal Difference learning

Branches of Machine Learning




Reinforcement Learning



Applications



Cooling and resource management systems



Source: DeepMind


Applications

Nuclear fusion control



[https://www.deepmind.com/blog/article/
Accelerating-fusion-science-through-learned-plasma-control](https://www.deepmind.com/blog/article/Accelerating-fusion-science-through-learned-plasma-control)

Multi-armed Bandit



Formally

There are k discrete *actions*: $a \in \mathcal{A}$

There is a **reward** associated with each action: $R(a)$
this is a random variable

Goal: give actions a_1, a_2, \dots so that the **cumulative reward** is maximal:


$$\arg \max_{a_1, a_2, \dots} \mathbb{E} \left\{ \sum_{t=1}^{\infty} R(a_t) \right\}$$

Catch: you don't know R !

You can only observe the realizations

Fun Time

Which action would you choose?



Reference answer: 1

Demo for home: [https:](https://iosband.github.io/2015/07/28/Beat-the-bandit.html)

[//iosband.github.io/2015/07/28/Beat-the-bandit.html](https://iosband.github.io/2015/07/28/Beat-the-bandit.html)

How to solve this?

Value estimation

Define the **value** of an action:

$$Q(a) = \mathbb{E}\{R(a)\}$$

Estimate it by averaging:

$$Q(a) \approx q(a) = \frac{1}{T} \sum_{t=1}^T r_t,$$


where r_t are the concrete realisations (sampled according to the distribution of $R(a)$).

Then simply **choose** $\arg \max_a q(a)$.

That's what you call interaction?

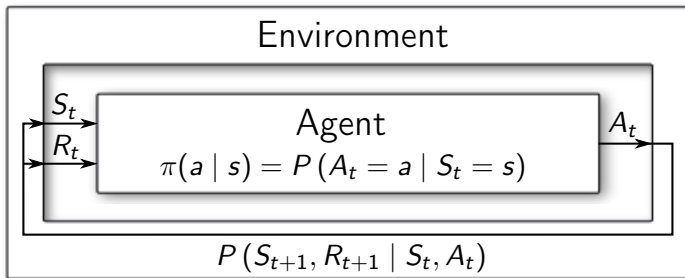


No, this is interaction



General formulation

with a figure



Markov-property: next state depends only on current state
 \implies probability of the whole trajectory is a product of probs. of transitions

Markov Decision Process

Let's say we have an environment with **states** $s \in \mathcal{S}$ and possible **actions** $a \in \mathcal{A}$.

The **dynamics** \mathcal{P} gives the probabilities of the next state s_{t+1} given the current state s_t and the given action a_t .

The **reward function** $R(s, a)$ gives the expected reward for action a given in state s .

The **policy** $\pi(a | s)$ of the agent is the probability of taking action a in state s .

Abuse of notation: *deterministic policy* is denoted by $a = \pi(s)$ (ie. taking action a w. p. 1).


Expected (undiscounted) **return**:

$$G^\pi(s) = \mathbb{E}_{a_t \sim \pi} \mathbb{E}_{s_t \sim \mathcal{P}} \left\{ \sum_{t=0}^T R(s_t, a_t) \right\}$$

Goal: find a policy $\pi = \arg \max_{\pi} G^\pi(s_0)$ (for some initial state s_0).

Fun Time

What is the *minimal* expected return in the following environment?



- From **S** Start to **G** Goal
- Reward for each step is -0.01
- Reward for stepping onto water (blue): -0.1
- Reward for hitting a wall: -0.5
- Reward for stepping on goal: $+1$

① 0 ② 0.33 ③ 0.83 ④ 0.93 ⑤ 1

Reference answer: 4



$$7 \cdot (-0.01) + 1 = 0.93$$

Stochastic Gradient Ascent

Objective: maximise

$$G^\pi = \mathbb{E}_{a_t \sim \pi, s_t, r_t \sim \mathcal{P}} \left\{ \sum_{t=1}^T r_t \right\}$$

Idea: $\pi = \pi_\theta$ with some parameter θ , and do gradient ascent:

$$\theta \leftarrow \theta + \alpha \nabla_\theta G^\pi$$

Calculating the gradient

Let $\tau = (s_0, a_0, s_1, \dots, s_T)$ the trajectory.

$$\begin{aligned}\nabla_{\theta} G^{\pi} &= \nabla_{\theta} \mathbb{E}_{a_t \sim \pi_{\theta}} \mathbb{E}_{s_t \sim \mathcal{P}} R(\tau) \\ &= \mathbb{E}_{a_t \sim \pi_{\theta}} \mathbb{E}_{s_t \sim \mathcal{P}} \nabla_{\theta} \log P_{\theta}(\tau) R(\tau) \\ &= \mathbb{E}_{a_t} \mathbb{E}_{s_t} \left\{ \nabla_{\theta} \log \prod_t \pi_{\theta}(a_t \mid s_t) P(s_{t+1} \mid s_t, a_t) R(\tau) \right\} \\ &= \mathbb{E}_{a_t} \mathbb{E}_{s_t} \left\{ \nabla_{\theta} \left[\sum_t \log \pi_{\theta}(a_t \mid s_t) + \log P(s_{t+1} \mid s_t, a_t) \right] R(\tau) \right\} \\ &= \mathbb{E}_{a_t} \mathbb{E}_{s_t} \left\{ \left[\sum_t \nabla_{\theta} \log \pi_{\theta}(a_t \mid s_t) \right] R(\tau) \right\}\end{aligned}$$

A Policy Gradient algorithm


$$\theta \leftarrow \theta + \alpha \nabla_{\theta} G^{\pi}$$

$$\nabla_{\theta} G^{\phi} = \mathbb{E}_{a_t} \mathbb{E}_{s_t} \left\{ \left[\sum_t \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \right] r(\tau) \right\}$$


The REINFORCE algorithm:

- ➊ run your agent (π_{θ}), sample some trajectories $\{\tau^{(i)}\}$
- ➋ estimate $G^{\pi} \approx \frac{1}{N} \sum_i \hat{G}^{\pi}(\tau^{(i)})$ (by the return)
- ➌ estimate the gradient
- ➍ update: $\theta \leftarrow \theta + \alpha \nabla_{\theta} \hat{G}^{\pi}$
- ➎ repeat


What's wrong with PG?



What's wrong with PG?



What's wrong with PG?



Variance reduction

by exploiting causality

Problem: high variance of the gradient estimation

But the policy at time t can't affect reward at time $t' < t$:

$$\nabla_{\theta} G^{\phi} = \mathbb{E}_{a_t} \mathbb{E}_{s_t} \left\{ \left[\sum_t \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \left(\sum_{t'=\mathbf{t}}^T R(s_{t'}, a_{t'}) \right) \right] \right\}$$

Variance reduction

by subtracting a baseline

Have a look at this:

$$\mathbb{E}_{a_t} \mathbb{E}_{s_t} \left\{ \left[\sum_t \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \right] (r(\tau) - b) \right\}$$

for some $b \in \mathbb{R}$.

$$\mathbb{E}\{\nabla_{\theta} \log P_{\theta}(\tau) b\} = b \mathbb{E}\{\nabla_{\theta} \log P_{\theta}(\tau)\} = b \nabla_{\theta} \mathbb{E}\{1\} = b \nabla_{\theta} 1 = 0$$

So this is also an unbiased estimator, but the variance is different!

$b = \bar{R} = \frac{1}{N} \sum_i r_i$ works in practice and is easy to compute.

How good is a policy?

Remember the bandits? We wanted to choose the arm with the maximal expected reward:


$$\arg \max_a \mathbb{E}\{R(a)\} =: Q(a)$$

Similarly in MDPs:

$$Q(s, a) := \mathbb{E}\left\{\sum_{t=1}^T R(s_t, a_t) \mid s_0 = s, a_0 = a, \pi\right\}$$

the *value* of the state and action for a policy π

Graphical intuition



Temporal Difference (TD) learning: based on the difference of the value function in two consecutive timesteps.

Bellman equation

We want to estimate how well the policy performs:

$$Q(s, a) = R(s, a) + \mathbb{E}_{s' \sim \mathcal{P}} \mathbb{E}_{a' \sim \pi} \{Q(s', a')\}$$

($\forall s, a$) This is the [Bellman equation](#).

Well, let's update during interaction: at a specific step t :

$$Q(s_t, a_t) \leftarrow r_t + Q(s_{t+1}, \pi(s_{t+1}))$$

Do this many times, and Q will converge for a policy.

Policy improvement

So far, the policy π stayed fixed. How to improve?

Idea: Greedy policy wrt Q :

$$\pi(s) = \arg \max_a Q(s, a)$$

Note: *deterministic policy*

Provably not worse

Generalized Policy Iteration

Putting it together:

- estimate Q for a given π
- improve π
 $\pi(s) \leftarrow \arg \max_a Q(s, a)$
- repeat

A bit slow...

Idea #2: do these two in parallel: *Generalized Policy Iteration*

$$Q(s_t, a_t) \leftarrow r_t + \max_{a'} Q(s_{t+1}, a')$$

Q-learning

Q-learning (off-policy TD control) for estimating $\pi \approx \pi_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+$, $a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

 Initialize S

 Loop for each step of episode:

 Choose A from S using policy derived from Q (e.g., ε -greedy)

 Take action A , observe R, S'

$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$

$S \leftarrow S'$

 until S is terminal

Will provably converge under some assumptions.

Note that it (as is) works only when $|\mathcal{S}| < \infty$.

State of the art algorithms

- Policy Gradient → [Proximal Policy Optimization \(PPO\)](#)
 - ChatGPT fine-tuning, Dota
- Q learning → [Deep Q Network \(DQN\)](#) → [Rainbow](#) → [Agent57](#)
 - Atari games
- [Maximum a Posteriori Policy Optimization \(MPO\)](#)
 - “Suppose the agent will perform well, what is the probability for this action?”
 - AlphaTensor, plasma controller
- Monte Carlo Tree Search (MTCS) → [AlphaGo 0](#) → [MuZero](#)
 - Assumes known dynamics; similar to planning
 - Go, Chess, Atari games
- PG + Q learning → Actor Critic → [IMPALA](#)
 - AlphaStar

Further reading

- [1] R. S. Sutton and A. G. Barto, *Reinforcement learning: An introduction*, 2nd ed. MIT press, 2018. [Online]. Available: <http://incompleteideas.net/book/the-book-2nd.html> (visited on 05/12/2018).
- [2] C. Szepesvári, “Reinforcement learning algorithms for MDPs,” *Synthesis lectures on artificial intelligence and machine learning*, vol. 4, no. 1, pp. 1–103, 2010.
- [3] Y. Li, “Deep reinforcement learning,” Oct. 15, 2018. arXiv: <http://arxiv.org/abs/1810.06339v1> [cs.LG].
 - DRL course at UC Berkeley
 - Spinning Up in Deep RL
 - EEML RL tutorials: 2021, 2022