

Support Vector Machine and PCA

Tasks

1. Implement the SVM algorithm using the unconstrained optimisation equations from the lectures and the `scipy.optimize.minimize` method in Scipy, which is a black box optimiser.

For this, you will need to implement a function that returns the value of the objective given the current α vector. See more in the [documentation](#). Calculate the Q matrix beforehand, do not calculate it in every call of this function¹!

- a. Using the algorithm, classify the dataset found in `svm_linear.npz`. Use C values of 1, 10 and 100.
 - b. How many nonzero (accounting for rounding/computational errors) elements does the α vector have? I.e. how many support vectors did the algorithm find?
 - c. Plot the data and the separator line (as usual), and mark the support vectors on your plot. What are the scores (i.e. $y(w^T x + b)$) for these samples? (You can use the `visualise_nonlin` function.)
 - d. Print the α_i s for the support vectors on your plot.
 - e. Compare the result of the SVM with the PLA and the Logistic Regression. Which one had the best generalisation (in this case)?
2. (*Optional*) Using the kernel trick, classify the dataset found in `rings.npz` from Lab 4. I would suggest using a third degree polynomial kernel for this.

What is the test/train accuracy of the classification? Plot the data and the separator curve. (For the plot, you can use the `visualise_nonlin` function.)

Mark the support vectors again; how many of them are there?

3. Given the dataset `PCA.npz` that was generated by applying a linear transformation and some noise to a 2D dataset; apply the PCA algorithm to find the two-dimensional factors underlying the data!

The final orientation should be apparent from visualising (a scatterplot of) the 2D data². What was the original transformation matrix used to create the transformed dataset?

¹Note that you can only do this if you do not have large amount of data, because it is quadratic in N .

²Note: I am left-handed.