



Pázmány Péter Catholic University
Faculty of Information Technology and Bionics

Parallel Programming

Lecture 2

István Reguly

reguly.istvan@itk.ppke.hu



Pázmány Péter Catholic University
Faculty of Information Technology and Bionics

So how do we parallelise an application?



Understand the problem

- Find the hotspots/bottlenecks
 - Know where most of the time is spent, where real work is done
 - Identify the underlying algorithm, and try to parallelise it
- Can it even be parallelised?
 - Yes: calculate the potential energy for each of 10 thousand independent conformations of a molecule. When done, find the minimum energy conformation.
 - No: calculate the Fibonacci series by using the formula $F(n) = F(n-1) + F(n-2)$




Typical steps

- Identify what pieces of work can be done concurrently
- Partition concurrent work onto independent processors
- Coordinate accesses to shared data (avoid conflicts)
- Ensure proper order of work using synchronisation
- What is “typical”
 - In some cases, some of these steps are unnecessary
 - Mapping of work to processors can be done either manually or automatically



Example: Dense matrix-vector product




- Computing each element of the output y is independent
- Easy to decompose into tasks: one per element in y
- Properties:
 - Task size is uniform
 - No dependence between tasks
 - Tasks share b (read)



Creating tasks

- Divide work into a number of tasks, identify how they are interrelated
- Many different decompositions possible
- Tasks may be the same, different or varying sizes
- Conceptualise with a Directed Acyclic Graph (DAG)





Example: Database Query

- Consider the execution of the query:

**MODEL = "CIVIC" AND YEAR=2001 AND
(COLOR="GREEN" OR COLOR="WHITE")**

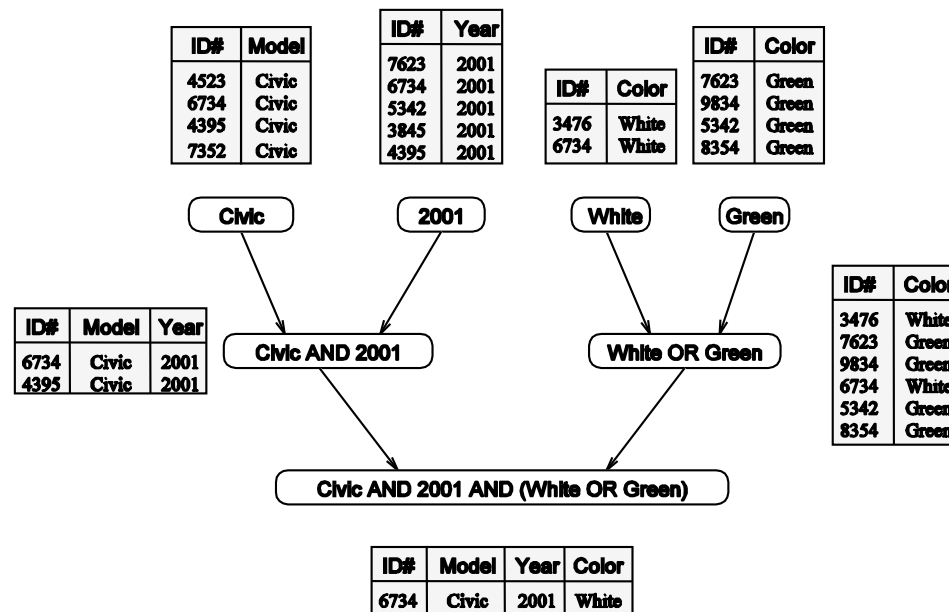
ID#	Model	Year	Color	Dealer	Price
4523	Civic	2002	Blue	MN	\$18,000
3476	Corolla	1999	White	IL	\$15,000
7623	Camry	2001	Green	NY	\$21,000
9834	Prius	2001	Green	CA	\$18,000
6734	Civic	2001	White	OR	\$17,000
5342	Altima	2001	Green	FL	\$19,000
3845	Maxima	2001	Blue	NY	\$22,000
8354	Accord	2000	Green	VT	\$18,000
4395	Civic	2001	Red	CA	\$17,000
7352	Civic	2002	Red	WA	\$18,000



Example: Database Query

- Task: select set of elements that satisfy a predicate
- Edge: output of one task is the input of the next

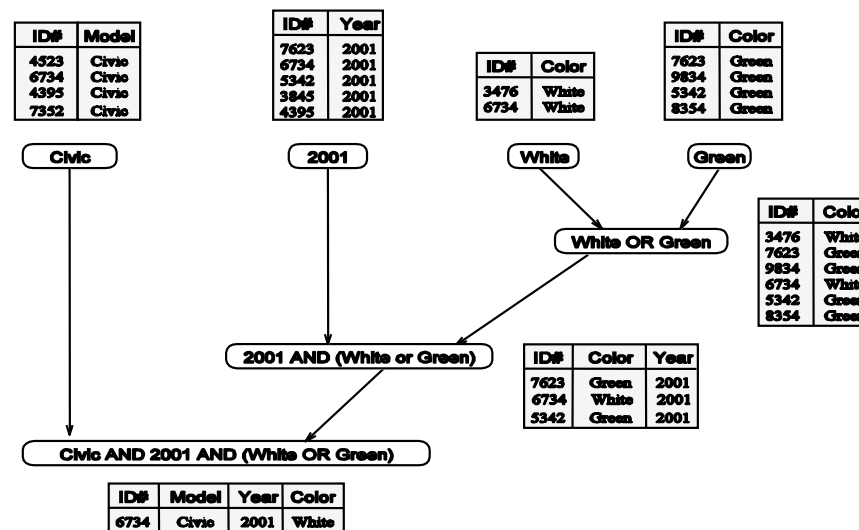
MODEL = "CIVIC" AND YEAR = 2001 AND
(COLOR = "GREEN" OR COLOR = "WHITE")





Example: Database Query

MODEL = "CIVIC" AND YEAR = 2001 AND
(COLOR = "GREEN" OR COLOR = "WHITE")




Different decompositions can lead to different parallelism and different amounts of work



Task granularity

- Granularity = task size
- Fine-grain: large number of small tasks
- Coarse-grain: small number of large tasks
- For dense matrix-vector multiply:
 - Fine-grain: one task per element of y
 - Coarse-grain: one task per 3 elements of y






Degree of concurrency

- Definition: the number of tasks that can execute in parallel
- May change during execution
- Metrics:
 - Maximum degree of concurrency
 - Largest number of concurrent tasks at any point
 - Average degree of concurrency
 - Average number of tasks that can be processed in parallel
- Degree of concurrency vs. task granularity
 - Inverse relationship



Critical path

- Edge in task dependency graph represents task serialization
- Critical path = longest path through the graph
 - Represents a lower bound on parallel execution time



Number in vertex is cost



Thought experiment


- N people, how many handshakes?
 - In serial?
 - In parallel?
- Molecular dynamics: N atoms – what is the total force on each
 - Each has a position – compute pairwise forces, then compute displacement



Limits of Parallel programming

- Amdahl's law: the potential program speedup is defined by the fraction of work(P) that can be parallelised:
- $P=0$, speedup=1 (none)
- $P=1$, speedup infinite (in theory)

$$speedup = \frac{1}{1 - P}$$






Amdahl's Law

- Given N processors:
- Limits to the scalability!
- In most cases, some serial fraction is unavoidable

$$speedup = \frac{1}{\underbrace{(1 - P)}_{\text{Serial}} + \underbrace{P/N}_{\text{Parallel}}}$$

All is lost?





Scaling the problem

- In many cases, trying to solve a larger problem helps: the serial overhead is constant.
 - E.g. 2-factor PDE computation, for a given size:

2D calculation:	85 seconds	85%
Serial overhead:	15 seconds	15%

- Doubling the discretisation points in each dimension and 2x the timesteps:

2D calculation:	680 seconds	97.84%
Serial overhead:	15 seconds	2.16%

Leaves to Gustafson' Law...



Measuring parallel performance

- We have a number of theoretical metrics
 - Ratio of sequential and parallel parts (Amdahl's Law)
 - Average/Maximum degree of parallelism
 - Task granularity, etc...
 - But these only help with the analysis of parallel algorithms, and give some theoretical bounds – they do little in the way of predicting how fast our implementation will be
 - Furthermore, they are usually focused on computations, not data movement
- The most straightforward way to measure parallel performance is to compare the execution times of the sequential program to the parallel one
 - $\text{Speedup} = \text{sequential time} / \text{parallel time}$



Matrix-Matrix multiply OpenMP

```
#pragma omp parallel for
for (int i=0; i < N; i++)
    for (int j=0; j < N; j++)
        for (int k=0; k < N; k++)
            c[i*N+j] += a[i*N+k] * b[k*N+j];
```

- It's that simple with OpenMP
 - Task granularity is chosen automatically
 - But there is only N-way concurrency and N^3 operations – could this be improved somehow?




Exercise

- Implement a matrix-matrix multiplication
 - Complete matmat.cpp skeleton
 - Test the `-Ofast` and the `-O0` flags' effect
- Add `#pragma omp parallel` for in front of the outermost loop and compile with `-fopenmp`
 - Observe the speedups and overheads
- Compute theoretical number of operations
 - Compute operations/second
- Optional exercise: create a graph which shows achieved operations/second at different matrix sizes – 100 to 5000
 - And with/without OpenMP



Shared Memory Model

- Tasks (processes/threads) share a common address space, which they read/write to asynchronously
- Mechanisms to control access and prevent race conditions (locks, semaphores)
- No data “ownership” – equal access -> fairly simple
- **Data locality:** difficult to manage
 - Keeping data local to where we work on it speeds up accesses
 - Difficult to control





Shared memory programming

- Most parallel programming models work in a shared memory environment
 - All SIMD models are naturally shared memory
- Question is how do we control parallelism, synchronisation, data access to shared and private data
 - We tend to call concurrent streams of instructions that can address the same shared memory “threads”
 - Threads are spawned (created) by a master process, they have their own state (e.g. where in the instruction stream they are, and some private data)
 - Different programming models give different ways to control threads



Threads model

- Implementation:
 - From a programming perspective, we need a set of API calls from within the source code, as well as compiler directives in the code to indicate parallel execution of tasks
- OpenMP
 - Industry standard: jointly defined by a group of HW/SW organizations
 - Compiler directive based
 - Portable (OS and languages, C/Fortran)
 - Easy to use – can be added incrementally




OpenMP



- Open Multi-Processing
- Three components:
 - Compiler directives
 - Runtime library routines
 - Environment variables
- High-level model
 - Implicit mapping and load balancing of work
- Standard
- Portable
- For more: <https://computing.llnl.gov/tutorials/openMP/>




OpenMP view of memory



Uniform Memory Access


Reality may be:
Uniform Memory Access





Fork-join model

- All programs begin as a single process: the **master thread**. It executes sequentially until the first **parallel region** construct.
- FORK: master thread creates a team of parallel threads
- Statements within the parallel region are executed in parallel among the different threads
- JOIN: Having completed the parallel region, threads synchronize and terminate, leaving the master thread.






Fork-join model

```
#pragma omp parallel for          Fork  
for (int i=0; i < N; i++)  
    c[i] = a[i] + b[i];          Join  
std::cout << "step 1\n"
```

```
#pragma omp parallel for          Fork  
for (int i=0; i < N; i++)  
    d[i] = c[i] + b[i];          Join  
std::cout << "step 2\n"
```

```
#pragma omp parallel for          Fork  
for (int i=0; i < N; i++)  
    e[i] = d[i] + c[i];          Join  
std::cout << "step 3\n"
```





Compiler directives

- Appear as comments in the code and ignored by the compiler unless instructed to (compiler flag)
- OpenMP directives are used for various purposes:
 - Spawning a parallel region
 - Dividing blocks of code among threads
 - Distributing loop iterations between threads
 - Serializing sections of code
 - Synchronization of work among threads
- Syntax:

```
Sentinel    directive-name          [clause, ...]  
#pragma omp parallel default(shared) private(beta,pi)
```



Runtime Library Routines

- The OpenMP API includes a growing number of library routines
- Used for many things:
 - Setting and querying the number of threads
 - Querying the threads unique identifier
 - Setting/querying dynamic threads, nested parallelism
 - Querying time
 - etc...

```
#include <omp.h>  
int omp_get_num_threads(void)
```



Environment variables

- OpenMP provides several environment variables to control the execution of parallel code
- Can be used to:
 - Set number of threads
 - Specify how loop iterations are divided
 - Binding threads to physical processors
 - Controlling nested parallelism, dynamic threads
 - etc.

```
export OMP_NUM_THREADS=4
```



Directives

#pragma omp	Directive-name	[clause,...]	newline
Required for all directives	A valid OpenMP directive.	Optional. In any order and repeated as necessary	Required. Precedes the structured block which is enclosed

```
#pragma omp parallel default(shared) private(beta,pi)
```

Rules:

- Case sensitive
- Only one directive-name per directive (there are a few combinations)
- Each directive applies to at most one succeeding statement, which must be a structured block
- Long directive lines can be "continued" on succeeding lines with "\"



Parallel region construct

- A parallel region is a block of code that will be executed by multiple threads. This is the fundamental OpenMP parallel construct.
- Format:

```
#pragma omp parallel [clause ...] newline  
    if (scalar_expression)  
    private (list)  
    shared (list)  
    default (shared | none)  
    firstprivate (list)  
    reduction (operator: list)  
    copyin (list)  
    num_threads (integer-expression)  
  
    { ... structured_block... }
```

Optional parameters



Parallel region

- When a thread reaches a parallel directive, it creates a team of threads and becomes the master of the team. The master is a member, and has ID 0.
- Starting from the beginning of the region, the code is “duplicated” and all threads will execute that code
 - Variables defined inside a parallel region, will be replicated too
- There is an implied barrier (synchronisation point) at the end of the parallel section, and only the master continues execution beyond it.
- If a thread terminates within a parallel region, all threads in the team will terminate and the work done is undefined



```
#include <omp.h>
main () {
int var1, var2, var3;
```

Serial code

*Beginning of parallel section. Fork a team of threads.
Specify variable scoping*

```
#pragma omp parallel private(var1, var2) shared(var3)
{
Parallel section executed by all threads
Other OpenMP directives
Run-time Library calls
All threads join master thread and disband
}
Resume serial code
}
```



Scoping

```
int main() {
    #pragma omp parallel
    {
        ...
        myfun()
    }
}
void myfun()
{
    ...
    #pragma omp critical
    {
        ...
    }
    ...
    #pragma omp sections
    {
        ...
    }
    ...
}
```



Directive Scoping

- Static Extent:
 - The code textually enclosed between the beginning and the end of a structured block following a directive
 - The static extent of a directive does not span multiple routines or source files
- Orphaned directive:
 - An OpenMP directive that appears independently from another enclosing directive: exists outside of another directive's static extent
- Dynamic extent:
 - Includes both its static extent and the extents of its orphaned directives
- Why? – There are some scoping rules on two directives can associate and nest within each other



Scoping

```
int main() {  
    #pragma omp parallel  
    {  
        ...  
        myfun()  
    }  
}  
void myfun()  
...  
#pragma omp critical  
{  
    ...  
}  
...  
#pragma omp sections  
{  
    ...  
}  
...
```

Static Extent

Orphaned Directives

Dynamic Extent



How many threads?

- The number of threads depends on the following:
 1. The `if()` clause (has to be true, otherwise serial)
 2. The `num_threads()` clause
 3. Use of the `set_omp_num_threads()` API
 4. The `OMP_NUM_THREADS` environment variable
 5. Implementation default (number of CPUs)
- Numbering is from 0 (master) to N-1



Exercise - omp_hello.cpp

```
#include <omp.h>
#include <iostream>

int main() {
    int nthreads, tid;

    /* Fork a team of threads with each thread
       having a private tid variable */
    #pragma omp parallel private(tid)
    {

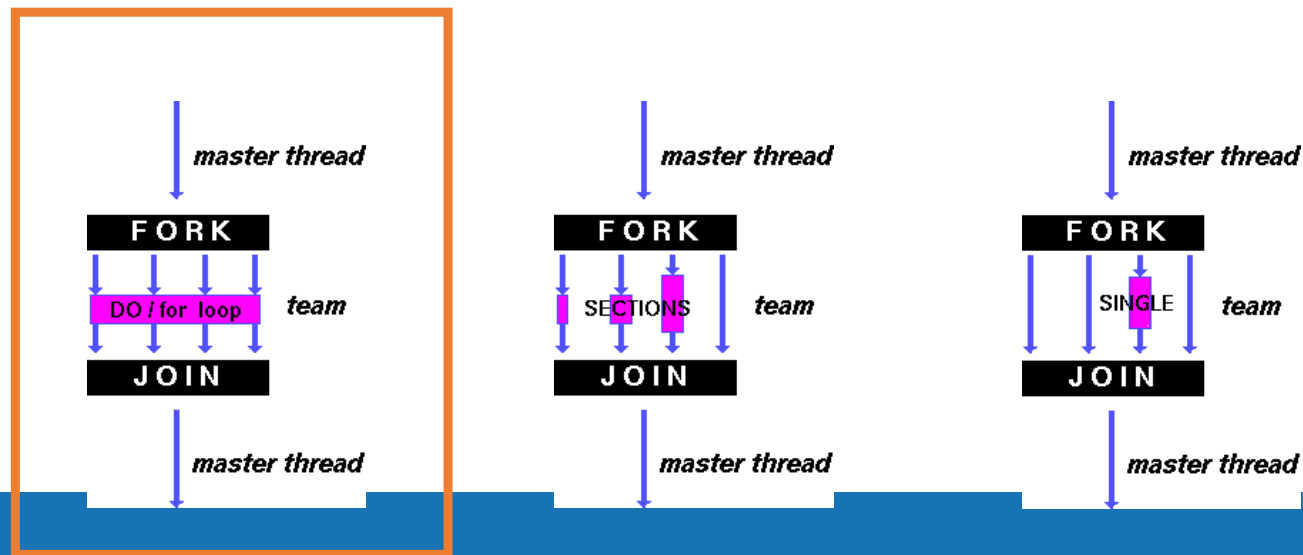
        /* Obtain and print thread id */
        tid = omp_get_thread_num();
        std::cout << "Hello World from thread = " << tid << std::endl;

        /* Only the master thread does this */
        if (tid == 0)
        {
            nthreads = omp_get_num_threads();
            std::cout << "Number of threads = " << nthreads << std::endl;
        }
    } /* All threads join master thread and terminate */
}
```



Work-sharing constructs

- Divides the execution of the enclosed code region among the members of the team that encounter it
- Does not launch new threads
- No implied barrier upon entry, but there is one at the exit.
- Has to occur within parallel region, encountered by all members, in the same order





for Directive

- The for directive specifies that the iterations of the loop immediately following it must be executed in parallel.

#pragma omp for *[clause ...] newline* Optional parameters

schedule (type [,chunk])
ordered
private (list)
firstprivate (list)
lastprivate (list)
shared (list)
reduction (operator: list)
collapse (n)
nowait

for_loop



Clauses

- **Schedule:** describes how iterations are divided among the threads. No default, the optimal one depends on the problem
 - **Static:** pieces of size *chunk*, statically assigned to threads. Default *chunk* is $\text{size}/\text{num_threads}$.
 - **Dynamic:** pieces of size *chunk*, dynamically scheduled. Default *chunk* is 1.
 - **Guided:** Like dynamic, but *chunk* size continuously decreases, starts out with $\text{size}/\text{num_threads}$, then $\text{leftover}/\text{num_threads}$, etc.
 - **Runtime:** deferred, depends on OMP_SCHEDULE environment variable.
 - **Auto:** let the compiler/runtime decide



Clauses

- **Nowait:** no synchronisation at the end of the loop
- **Ordered:** requires that some parts of the iterations are executed as in a serial program + `#pragma omp ordered`
- **Collapse:** specifies how many nested loops to be collapsed into one large iteration space and then divided up
- Restrictions: has to be for loop with integer iteration variable



Vector add

```
#include <omp.h>
#include <iostream>
#include <chrono>
#include <vector>
#define N 100000000
int main (int argc, const char **argv)
{
    std::vector<float> a(N), b(N), c(N);
    /* Some initialisation */
    for (int i=0; i < N; i++)
        a[i] = b[i] = i * 1.0;
    auto t1 = std::chrono::high_resolution_clock::now();
#pragma omp parallel for
    for (int i=0; i < N; i++)
        c[i] = a[i] + b[i];
    auto t2 = std::chrono::high_resolution_clock::now();
    std::cout << "took "
        << std::chrono::duration_cast<std::chrono::milliseconds>(t2-t1).count()
        << " milliseconds\n";
}
```



Combinations

- Two combined directives for convenience:
 - Parallel for

```
#pragma omp parallel for \  
    shared(a,b,c,chunk) private(i) \  
    schedule(static,chunk)  
for (i=0; i < n; i++)  
    c[i] = a[i] + b[i];
```



Need for Synchronisation

THREAD 1:

```
increment(x)
```

```
{
```

```
    x = x + 1;
```

```
}
```

THREAD 1:

```
10  LOAD A, (x address)
```

```
20  ADD A, 1
```

```
30  STORE A, (x address)
```

THREAD 2:

```
increment(x)
```

```
{
```

```
    x = x + 1;
```

```
}
```

THREAD 2:

```
10  LOAD B, (x address)
```

```
20  ADD B, 1
```

```
30  STORE B, (x address)
```

A possible sequence:

1. Thread 1 loads x into register A ->
2. Thread 2 loads x into register B ->
3. Thread 1 adds 1 to A ->
4. Thread 2 adds 1 to B ->
5. Thread 1 writes A to x ->
6. Thread 2 writes B to x



Master directive

- Specifies a region that is to be executed only by the master thread of the team.
- No implied barrier

```
#pragma omp master  newline  
    structured_block
```



Critical directive

- A region (section) of code that must be executed by only one thread at a time

```
#pragma omp critical [ name ] newline
```

```
structured_block
```

- If a thread is in the critical section, and another arrives, it will block until the first finishes
- Can be named, so multiple sections in code are treated as one – all unnamed sections are treated as the same section



Critical section

```
#include <omp.h>

main()
{

    int x;
    x = 0;

    #pragma omp parallel shared(x)
    {

        #pragma omp critical
        x = x + 1;

    } /* end of parallel section */
    printf(„%d\n“, x);
}
```

Can this possibly be faster?



Critical section

```
typedef struct {  
    int index;  
    int value;  
} Event;  
  
...  
#pragma omp parallel for  
for(int i = 0; i < DATA_SIZE; i++) {  
    if(data[i] > THRESHOLD) {  
        #pragma omp critical  
        {  
            rareEvents[eventCount].index = i;  
            rareEvents[eventCount].value = data[i];  
            eventCount++;  
        }  
    }  
}
```



Barrier directive

- Synchronises all threads in the team
- When the barrier directive is reached, a thread will wait until all other threads have reached that barrier. Then all threads resume executing code after the barrier
- All threads (or none) must encounter the barrier

```
#pragma omp barrier newline
```



Barriers

```
#pragma omp parallel shared(image)
{
    // Step 1: Apply filter to assigned section
    applyFilter(image_section);

    // Barrier: Wait for all threads to finish Step 1
    #pragma omp barrier

    // Step 2: Apply the second enhancement
    applyEnhancement(image_section);
}
```



Atomic directive

- Specifies that a specific memory location must be updated atomically, rather than letting multiple threads attempt to write to it.
- It's like a mini critical section, applies to only one (following) statement

```
#pragma omp atomic newline  
statement_expression
```



Atomic directive

```
#pragma omp parallel for
for(int i = 0; i < DATA_SIZE; i++) {
    int bin = data[i] * NUM_BINS / MAX_VALUE; // Simple bin calculation
    #pragma omp atomic
    histogram[bin]++;
}
```



Data sharing

- Important to understand how variables are shared and how to use data scoping
- By default most variables are shared
 - Except loop index variables and anything created on the stack within a parallel region (i.e. local variables)
- Data scope attributes:
 - Private, firstprivate, lastprivate, shared, default, reduction, copyin
- Combined with other directives to control the scoping of enclosed variables



Private

- Lists variables that are to be private to each thread
- Behaviour:
 - A new object of the same type is declared once for each thread in the team
 - All references to the original are replaced with references to the new one
 - Should be assumed **uninitialised**
- Firstprivate combines private and initialisation
- Lastprivate writes back the value from the last iteration or section

```
private (list)
firstprivate (list)
lastprivate (list)
```



Shared

- Declares listed variables to be shared among all the threads in the team
- Exists only in one memory location, and all threads can read or write to that address
- Programmer's responsibility to ensure safe access

`shared (list)`



Reduction

- Performs a reduction on the variables in the list
- A private copy is created for each thread, and at the end they are reduced to the globally shared variable

```
reduction (operator: list)

#pragma omp parallel for      \
  default(shared) private(i) \
  reduction(+:result)

for (i=0; i < n; i++)
  result = result + (a[i] * b[i]);

printf("Final result= %f\n",result);
```



Reduction

- Variables have to be scalar (no arrays or structures)
- Have to be declared shared in any enclosing region
- Assumes associativity!

```
x = x op expr  
x = expr op x (except subtraction)  
x binop = expr  
x++  
++x  
x--  
--x
```

x is a scalar variable in the list

expr is a scalar expression that does not reference *x*


op is not overloaded, and is one of +, *, -, /, &, ^, |, &&, ||

binop is not overloaded, and is one of +, *, -, /, &, ^, |



Exercise - Compute Pi

- Open `calc_pi.cpp`
- Compile:
 - `g++ -std=c++11 calc_pi.cpp -o calc_pi -Ofast -fopenmp`
- Generate random points between `[-0.5 0.5]`
- Test if inside the circle
- Ratio of circle to square
- $\pi \approx 4 * (\text{\#inside} / \text{\#total})$
- `double pi = 4.0 * (inside / 10000.0);`





Semester project

- Computational fluid dynamics problem
- <https://classroom.github.com/a/jmyW7aJf>
- You can plot the output file with Matlab to see fluid flow



Assignment – due March 17 midnight

- Add timing to your code – time the “main time iteration”, print elapsed time at the end
- Parallelise the code with OpenMP
 - Reductions!
 - Make sure your code validates – gives (almost) the same velocities with/without OpenMP
- Push your changes to the GitHub Classroom repo – make sure it's there by looking at it on the web interface